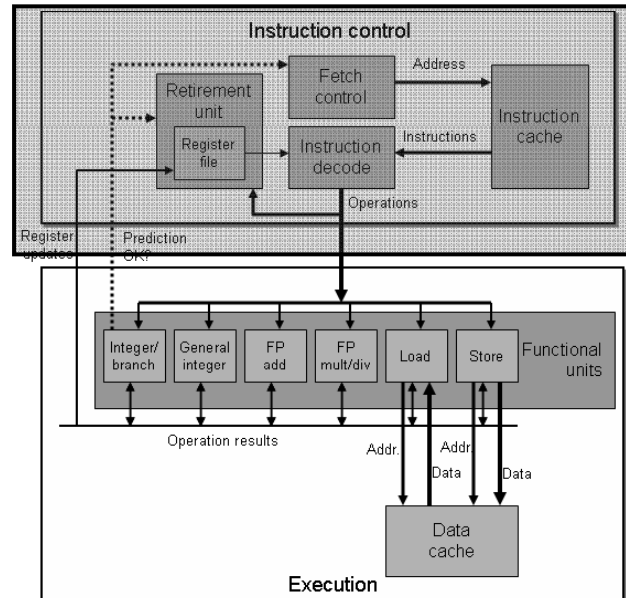


CSC-365
Computer Architecture

Chapter 5
Optimizing Program Performance
Section 5.7

Recall our Multiple Functional Units



Recall the Execution Unit (EU) Capabilities

- Multiple Instructions Can Execute in Parallel
 - 1 load
 - 1 store
 - 2 integer (one may be branch)
 - 1 FP Addition
 - 1 FP Multiplication or Division
- Benefit of multiple functional units
 - Since the Load/Store unit performs the MEM functions
 - It can now run in parallel to any EXE function
 - It also has its own Adder so we don't need the other functional units for it to do its job

Recall the Execution Unit (EU) Latencies

- Possible Instructions:

– Instruction	Latency	Issue Time
– Compare	1	1
– Integer Add	1	1
– Jump	1	1
– Load / Store	3	1
– Integer Multiply	4	1
– Integer Divide	36	36
– Double/Single FP Multiply	5	2
– Double/Single FP Add	3	1
– Double/Single FP Divide	38	38

Example of the Processing of the ICU (A loop to multiply all array elements)

- Code to multiply all the elements of an array together:
 - Integer data, multiply operation

```
.L24:                # Loop:
    imull (%eax,%edx,4),%ecx # t *= data[i]
    incl %edx                # i++
    cmpl %esi,%edx           # i:length
    jl .L24                  # if < goto Loop
```

- Translation of First Iteration

```
.L24:
    imull (%eax,%edx,4),%ecx

    incl %edx
    cmpl %esi,%edx
    jl .L24
```

```
load (%eax,%edx.0,4) → t.1
imull t.1, %ecx.0     → %ecx.1
incl %edx.0          → %edx.1
cmpl %esi, %edx.1     → cc.1
jl-taken cc.1
```

Recall the Use of Renaming for all Registers

- Instruction was originally:

`imull (%eax,%edx,4),%ecx`

- Which performed:

`%ecx = (%eax,%edx,4) * %ecx`

- We now avoid WAW and WAR hazards by replacing this with:

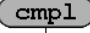
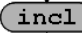
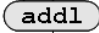


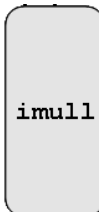
`load (%eax,%edx,4) → t.1`

`imull t.1, %ecx.0 → %ecx.1`

Note there are multiple copies of registers possible

Recall the Computational Graph Elements

- Possible Instructions:

– Instruction	Latency	Issue Time	Notation
– Compare	1	1	
– Integer Add/Increment	1	1	 
– Jump	1	1	
– Load / Store	3	1	 
– Integer Multiply	4	1	
– Integer Divide	36	36	
– Double/Single FP Multiply	5	2	
– Double/Single FP Add	3	1	
– Double/Single FP Divide	38	38	

Visualizing Operations in the EXE Stage for a The Following Instructions

- Recall the code to perform multiplications of all the elements of the array together
- It consists of two sets of operations
 - Actual multiplications of the array elements
 - Incrementing the index for the array

```

load (%eax,%edx,4) → t.1
imull t.1, %ecx.0 → %ecx.1
incl %edx.0 → %edx.1
cml %esi, %edx.1 → cc.1
jl-taken cc.1

```

First we can look at the multiplications in the instruction stream

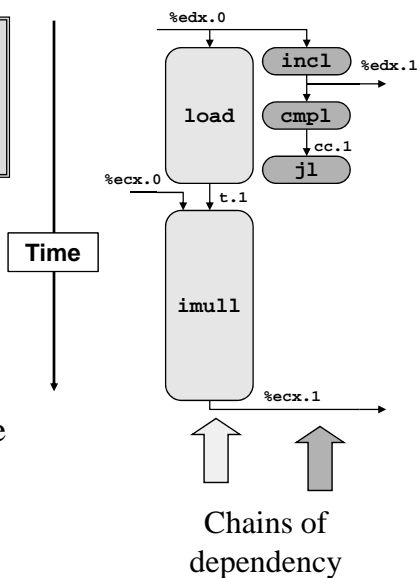
Recall Computational Graph in the EXE Stage for the Complete Array Multiplication

```

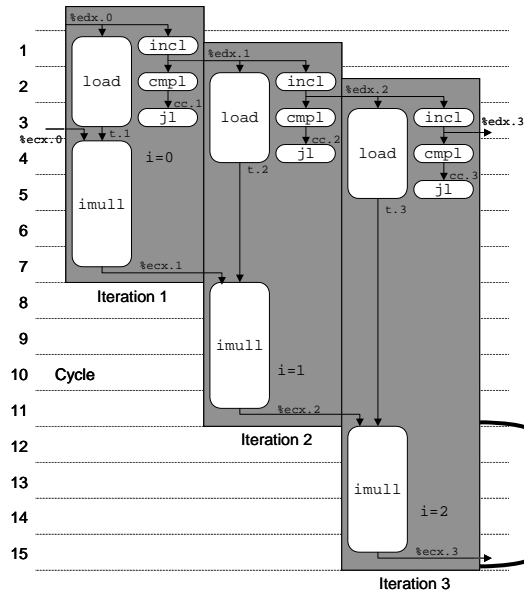
load (%eax,%edx,4) → t.1
imull t.1, %ecx.0 → %ecx.1
incl %edx.0 → %edx.1
cml %esi, %edx.1 → cc.1
jl-taken cc.1

```

- We can parallelize these instructions
 - Creates two unique paths of dependencies now
- Multiple lateral entries denote simultaneous instruction issues




3 Iterations of Combining Product



- Performance
 - Limiting factor becomes *latency* of integer multiplier
- Gives CPE of 4.0
 - Every 4 clock cycles a new result is generated

Benefits of Using the Computational Graph

- We can easily see the parallelism that is going on in our machine
- We can also see we are using Out-of-Order Execution



Assembly code has it here

```
load (%eax,%edx,4) → t.1
imull t.1, %ecx.0 → %ecx.1
incl %edx.0 → %edx.1
cpl %esi, %edx.1 → cc.1
j1-taken cc.1
```

- Instructions are being executed in the order in time that the resources are available
 - This order is different than when the assembly code dictated
 - Of course there must be no data dependencies between instructions whose order is changed

4 Iterations of Combining Product

4 integer ops

- Resources
 - Note we are using 4 integer units per cycle
 - We will need to fix this for finite resources

- Resources
 - Note we are using 4 integer units per cycle
 - We will need to fix this for finite resources

4 Iterations of Combining Product with 3 Integer ALU Resources

The diagram illustrates the execution of a loop over 4 iterations, utilizing 3 integer ALU resources. The vertical axis represents the cycle number (1 to 19), and the horizontal axis represents the instruction stream. A dashed line at cycle 6 indicates a point in time where the first three iterations have completed their integer operations, and the fourth iteration is beginning. The operations are scheduled as follows:

- Iteration 1:** `imul` (cycle 6), `incl` (cycle 7), `load` (cycle 8), `cmpl` (cycle 9), `jz` (cycle 10).
- Iteration 2:** `imul` (cycle 11), `incl` (cycle 12), `load` (cycle 13), `cmpl` (cycle 14), `jz` (cycle 15).
- Iteration 3:** `imul` (cycle 16), `incl` (cycle 17), `load` (cycle 18), `cmpl` (cycle 19), `jz` (cycle 20).
- Iteration 4:** `imul` (cycle 21), `incl` (cycle 22), `load` (cycle 23), `cmpl` (cycle 24), `jz` (cycle 25).

The diagram shows that the `imul` instruction is the most resource-intensive, requiring 3 ALU resources. The `incl`, `load`, `cmpl`, and `jz` instructions each require only 1 ALU resource. The overall execution time is 25 cycles, and the CPE (Cycles Per Element) is 4.

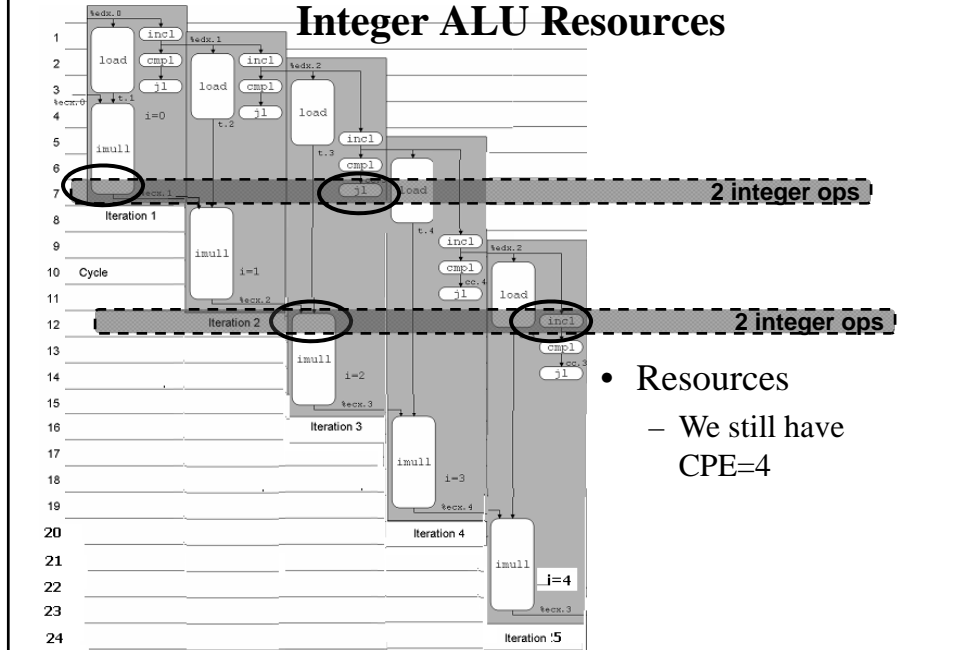
3 integer ops

Resources

- We now schedule around the available hardware
- Note it is 3 ops, but the mult instruction is now in the later pipe stages
- We still have CPE=4

- ## Resources
- We now schedule around the available hardware
 - Note it is 3 ops, but the mult instruction is now in the later pipe stages
 - We still have CPE=4

4 Iterations of Combining Product with 2 Integer ALU Resources



Operations in the EXE stage for Mult

- Notice that the CPE of the Multiplication example matches the latency of the Mult operation
 - We are limited by the amount of time the mult takes, even with an infinite amount of hardware
- Now let's look at an example where we replace the Mult operation with an Add
 - In other words we add all the elements of an array together
 - Recall the Add has a latency of only 1

Add Example for Hardware Scheduling

- Code to ADD all the elements of an array together:
 - Integer data, add operation

```
.L24:                # Loop:
    iaddl (%eax,%edx,4),%ecx # t += data[i]
    incl %edx                # i++
    cmpl %esi,%edx           # i:length
    jl .L24                  # if < goto Loop
```


- Translation of First Iteration

```
.L24:
    iaddl (%eax,%edx,4),%ecx

    incl %edx
    cmpl %esi,%edx
    jl .L24
```

```
load (%eax,%edx.0,4) → t.1
iaddl t.1, %ecx.0     → %ecx.1
incl %edx.0           → %edx.1
cmpl %esi, %edx.1     → cc.1
jl-taken cc.1
```


Visualizing Operations in the EXE stage for ADD



```
load (%eax,%edx,4) → t.1
iaddl t.1, %ecx.0 → %ecx.1
incl %edx.0       → %edx.1
cmpl %esi, %edx.1 → cc.1
jl-taken cc.1
```

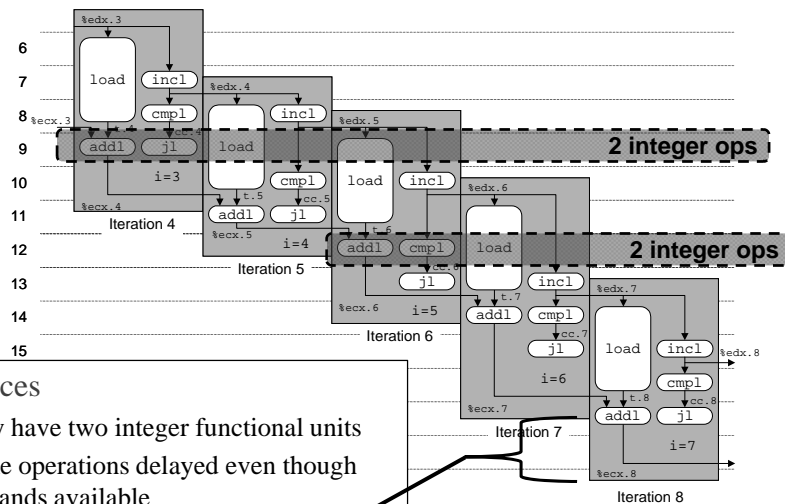
- Operations
 - Same as before, except that add has latency of 1
 - This simple change will have a dramatic impact on the scheduling

4 Iterations of Combining Sum



- Unlimited Resource Analysis
- Performance
 - Can begin a new iteration on each clock cycle
 - Should give CPE of 1.0
 - Would again require executing 4 integer operations in parallel !!!

Combining Sum: Resource Constraints



- Resources
 - Only have two integer functional units
 - Some operations delayed even though operands available
 - Set priority based on program order
- Performance
 - Sustain CPE of 2.0

Combining Sum: Resource Constraints II

- Note we now reschedule our instructions to match the availability of the hardware
- How do we actually schedule?
 - The details of this are extremely complicated
 - They are a major portion of CSC-565
- The general way is:
 - We track - what is the true order of execution of the instructions
 - Instruction priority is set based on its location in the machine code instruction streams
 - Earlier instructions are given higher priority than later instructions

Combining Sum: Resource Constraints III

- What do we see:
 - The performance of our system is now limited not by our usual data and control dependencies but now by hardware resources.
- There are now three limitations to performance:
 1. Data dependencies (Data Hazards)
 2. Control dependencies (Control Hazards)
 3. Hardware resource limitations
 - Called Structural Hazards

Combining Sum: Resource Constraints IV

- What to do?
 - Note that the first one is a limitation of our code
 - The second limitation we greatly reduced thru branch prediction
 - Is there anything we can do to minimize the impact of the third limitation?

Combining Sum: Resource Constraints V

- What limited our performance
 - There were not enough integer units to perform our add/mult operations or the incl, cmpl, or and jl instructions
 - Note the add and mult instructions were the only ones really involved with our calculation
 - The other three instructions were merely overhead for our loop that we used to compute the running product/sum

What to do??

- Therefore we need to find a way to reduce our loop overhead so we can free-up the functional units to perform the useful instructions
 - We do this with *Loop Unrolling*

Next Class

- Next class:
 - Loop unrolling Section 5.8
 - Begin Chapter 6